

# Apache Metron Profiler

**Nick Allen, Committer and PMC Member, Apache Metron**

DataWorks Summit / Hadoop Summit 2017

Birds of a Feather - Cyber Security and Apache Metron

Metron Version 0.4.0

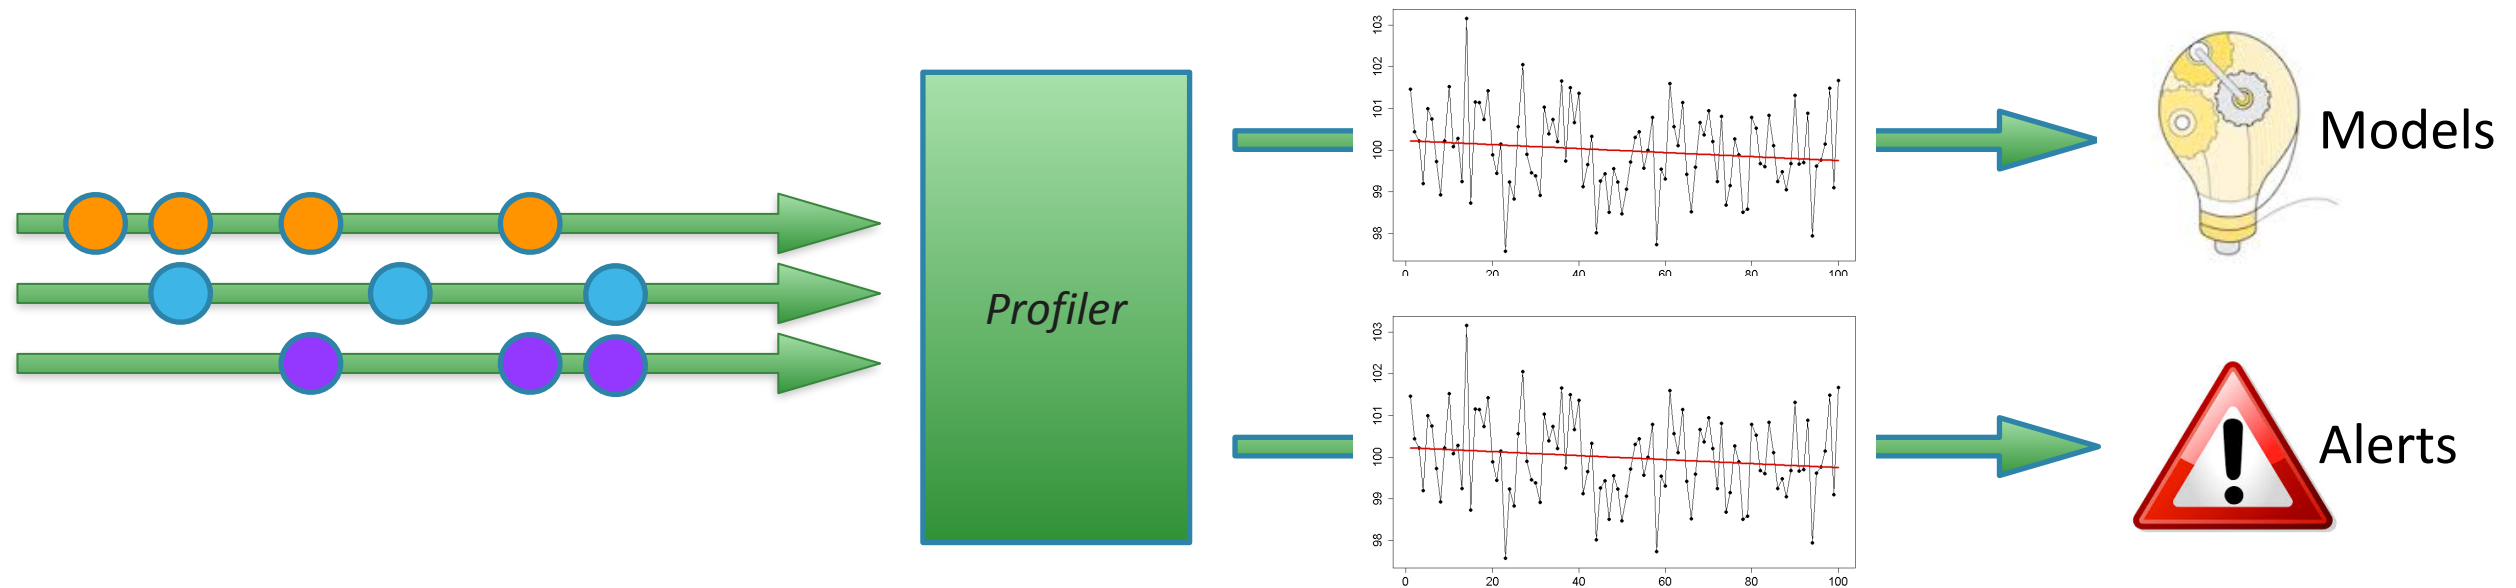


# Agenda

- ◆ Introduction
- ◆ “Hello World” Profile
- ◆ Data Sketches
- ◆ Profiles
- ◆ Implementation


# Introduction

# The Profiler




- ◆ A generalized, extensible solution for extracting feature sets from high throughput, streaming data
- ◆ Generates a profile describing the behavior of an entity; a host, user, subnet or application
- ◆ A foundational component for both security model building and alerting in Metron

## Background: Enrichments




- ◆ Enriching telemetry with contextual clues is invaluable
  - Dramatically improves threat triage and response
  - Another foundational component for security model building (expands the feature set)
- ◆ Security telemetry in Metron is 'sticky'
  - Enrichments 'stick' as the telemetry progresses through the system

# The Problem



- ◆ Enrichments operate within the context of a single message
  - Simple, efficient and scalable for most enrichment and triage scenarios
- ◆ Insufficient for other scenarios
  - Looking across time; trending
  - Looking across data sources; correlation
  - Looking at aggregate behaviors; How is an “application” or “user” behaving as a whole?

# The Profiler



- ◆ The Profiler creates logical windows that span both time and data sources
- ◆ The user defines a Profile that operates on these logical windows
  - A Profile consumes each message within a given window
  - Identifies an Entity; the subject of interest; a server, user, subnet, or application name
  - Produces a Result; a summary of an Entity's behavior within the window



# Profiles are Time Series

- A Profile executed on a given window results in a unique Result for each Entity


$$f(\text{window}_{t1}) \rightarrow (\text{entity}_1, \text{result}_{e1}), (\text{entity}_n, \text{result}_{en})$$

- As a Profile is executed over time, a series of Results starts to emerge


$$\begin{aligned} f(\text{window}_{t1}) &\rightarrow (\text{entity}_1, \text{result}_{t1}), (\text{entity}_n, \text{result}_{t1}) \\ f(\text{window}_{t2}) &\rightarrow (\text{entity}_1, \text{result}_{t2}), (\text{entity}_n, \text{result}_{t2}) \\ f(\text{window}_{tn}) &\rightarrow (\text{entity}_1, \text{result}_{tn}), (\text{entity}_n, \text{result}_{tn}) \end{aligned}$$

- Rearranging the data slightly, we can see that a unique time series results for each (Profile, Entity) pair

$$(\text{entity}_1) \rightarrow \text{result}_{t1}, \text{result}_{t2}, \text{result}_{tn}$$



$$(\text{entity}_n) \rightarrow \text{result}_{t1}, \text{result}_{t2}, \text{result}_{tn}$$





# Hello World

# “Hello World” Profile

- A Profile that counts the number of telemetry messages for each IP source address

```
{  
  "profile": "hello-world",  
  "foreach": "ip_src_addr",  
  "init":    { "count": "0" },  
  "update":  { "count": "count + 1" },  
  "result":  "count"  
}
```

# “Hello World” Profile

- ◆ A Profile that counts the number of telemetry messages for each IP source address

```
{  
  "profile": "hello-world",  
  "foreach": "ip_src_addr",  
  "init":    { "count": "0" },  
  "update":  { "count": "count + 1" },  
  "result":  "count"  
}
```

← Name; identifier for the profile

# “Hello World” Profile

- ◆ A Profile that counts the number of telemetry messages for each IP source address

```
{  
  "profile": "hello-world",  
  "foreach": "ip_src_addr",  
  "init":    { "count": "0" },  
  "update":  { "count": "count + 1" },  
  "result":  "count"  
}
```

← Name; identifier for the profile

← Entity; a Stellar expression  
Maintain a count for each unique source IP address

# “Hello World” Profile

- ◆ A Profile that counts the number of telemetry messages for each IP source address

```
{  
  "profile": "hello-world",  
  "foreach": "ip_src_addr",  
  "init":    { "count": "0" },  
  "update":  { "count": "count + 1" },  
  "result":  "count"  
}
```

← Name; identifier for the profile

← Entity; a Stellar expression  
Maintain a count for each unique source IP address

← Initialize a counter at the start of each window

# "Hello World" Profile

- ◆ A Profile that counts the number of telemetry messages for each IP source address

```
{  
  "profile": "hello-world",  
  "foreach": "ip_src_addr",  
  "init":    { "count": "0" },  
  "update":  { "count": "count + 1" },  
  "result":  "count"  
}
```

← Name; identifier for the profile

← Entity; a Stellar expression  
Maintain a count for each unique source IP address

← Initialize a counter at the start of each window

← Increment the counter for each message in the window

# "Hello World" Profile

- ◆ A Profile that counts the number of telemetry messages for each IP source address

```
{  
  "profile": "hello-world",  
  "foreach": "ip_src_addr",  
  "init":    { "count": "0" },  
  "update":  { "count": "count + 1" },  
  "result":  "count",  
}
```

← Name; identifier for the profile

← Entity; a Stellar expression  
Maintain a count for each unique source IP address

← Initialize a counter at the start of each window

← Increment the counter for each message in the window

← Return the count at the end of the window




# "Hello World" Profile

- ◆ A Profile that counts the number of telemetry messages for each IP source address

```
{  
  "profile": "hello-world",  
  "foreach": "ip_src_addr",  
  "init": { "count": "0" },  
  "update": { "count": "count + 1" },  
  "result": "count"  
}
```

Annotations:

- Name; identifier for the profile (points to "hello-world")
- Entity; a Stellar expression (points to "ip\_src\_addr")
- Maintain a count for each unique source IP address (points to "foreach")
- Initialize a counter at the start of each window (points to "count": "0")
- Increment the counter for each message in the window (points to "count + 1")
- Return the count at the end of the window (points to "count")



# “Hello World” Profile

- ◆ A Profile that counts the number of telemetry messages for each IP source address

```
{  
  "profile": "hello-world",  
  "foreach": "ip_src_addr",  
  "init": { "count": "0" },  
  "update": { "count": "count + 1" },  
  "result": "count"  
}
```

← Name; identifier for the profile

← Entity; a Stellar expression  
Maintain a count for each unique source IP address

← Initialize a counter at the start of each window

← Increment the counter for each message in the window


← Return the count at the end of the window

```
[Stellar]>>> PROFILE_GET("hello-world", "10.0.0.3", PROFILE_FIXED(30, "MINUTES"))  
[451, 448]
```

```
[Stellar]>>> PROFILE_GET("hello-world", "10.0.0.5", PROFILE_FIXED(30, "MINUTES"))  
[234, 176]
```

# Data Sketches

# Data Sketches



- ◆ Data Sketches provide fast, approximate answers to queries about the underlying data
- ◆ There are a variety of different types of data sketches, but general characteristics include
  - ◆ Stream Friendly - Each item of a stream, examined only once, can quickly update a small sketch data structure
  - ◆ Scalable - Effective for queries that do not scale well; count distinct, quantiles, most frequent items
  - ◆ Approximate, but with predictable error rates
  - ◆ Sub-Linear in Size - Required storage space grows more slowly than the input size
  - ◆ Mergeable (additive) and thus easily support parallelization
    - ◆  $\text{query}(\text{sketch}(\text{data1} + \text{data2})) == \text{query}(\text{sketch}(\text{data1}) + \text{sketch}(\text{data2}))$

# Data Sketches and the Profiler

- ◆ The Profiler can persist anything serializable; not just numbers
- ◆ Profiler + Data Sketches
  - ◆ Allows the Profiler to scale to very large data sets
  - ◆ Allows consumers to ask different queries of the same profile data
  - ◆ Allows consumers to change the time horizon
  - ◆ Produce small, lightweight objects that can be stored efficiently

# Data Sketches - Example

- ◆ A simple Profile that tracks URL length over time

```
{
  "profile": "http-length",
  "foreach": "'global'",
  "onlyif": "source.type == 'bro' and protocol == 'HTTP'",
  "update": { "sk": "STATS_ADD(sk, length)" },
  "result": "sk"
}
```

- ◆ These aren't just numbers

```
[Stellar]>>> stats := PROFILE_GET( "http-length", "global", PROFILE_FIXED(24, "HOURS"))
[Stellar]>>> stats
[org.apache.metron.common.math.stats.OnlineStatisticsProvider@79fe4ab9, ...]
```

- ◆ Ask different queries of the same data

```
[Stellar]>>> STATS_MEAN( GET_FIRST( stats))
15979.0625
[Stellar]>>> STATS_PERCENTILE( GET_FIRST(stats), 90)
30310.958
```

- ◆ Merge to change the time horizon

```
[Stellar]>>> merged := STATS_MERGE( stats)
[Stellar]>>> STATS_PERCENTILE(merged, 90)
29810.992
```

# Profiles



# Failed Logins




- ◆ A profile that tracks the number of bad logins by host

```
{  
  "profile": "bad-logins",  
  "foreach": "ip_src_addr",  
  "onlyif": "source.type == 'activedirectory' and event.type == 'failed_login'",  
  "init" : { "count" : "0" }  
  "update": { "count": "count + 1" },  
  "result": "count"  
}
```

# Vertex Degree

- View network communication as a directed graph
  - The in and out degree can distinguish behaviors
  - Anomalies over time can serve as an indicator of compromise

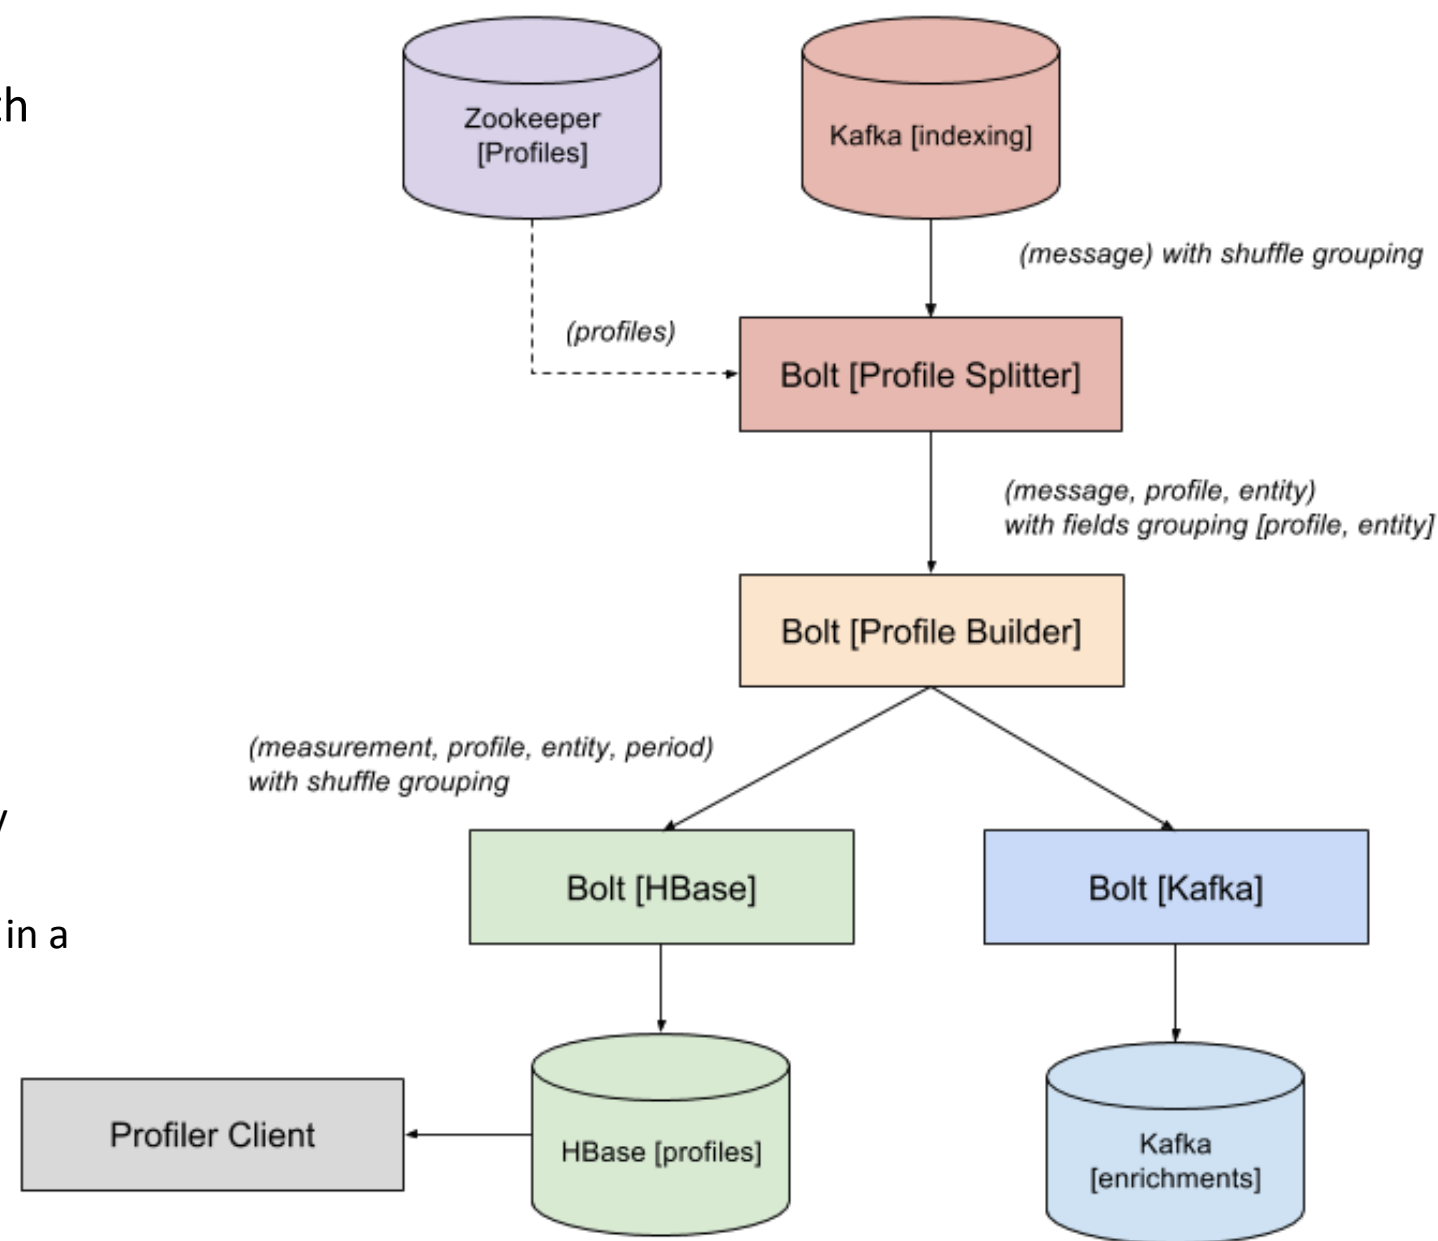
```
{  
  "profile": "in-degree",  
  "onlyif": "source.type == 'yaf'",  
  "foreach": "ip_dst_addr",  
  "init": { "in": "HLLP_INIT(5, 6)" },  
  "update": { "in": "HLLP_ADD(in, ip_src_addr)" },  
  "result": { "HLLP_CARDINALITY(in)" },  
}
```



# Implementation

# The Implementation

- ◆ A Storm topology that lives outside the critical path
- ◆ Profiles are defined in Zookeeper
- ◆ Profile Splitter Bolt
  - Reads all profile definitions in Zookeeper
  - Consumes each message and for each profile determines...
    - Is the message needed by the profile?
    - If needed, what is the entity?
  - Partitions the data by (Profile, Entity) using a fields grouping
- ◆ Profile Builder Bolt
  - Consumes the message, profile definition, and entity
  - Updates the state for a (Profile, Entity) pair
  - Flushes all state upon receiving a tick tuple resulting in a Profile Measurement



# The Implementation

## ■ HBase Bolt

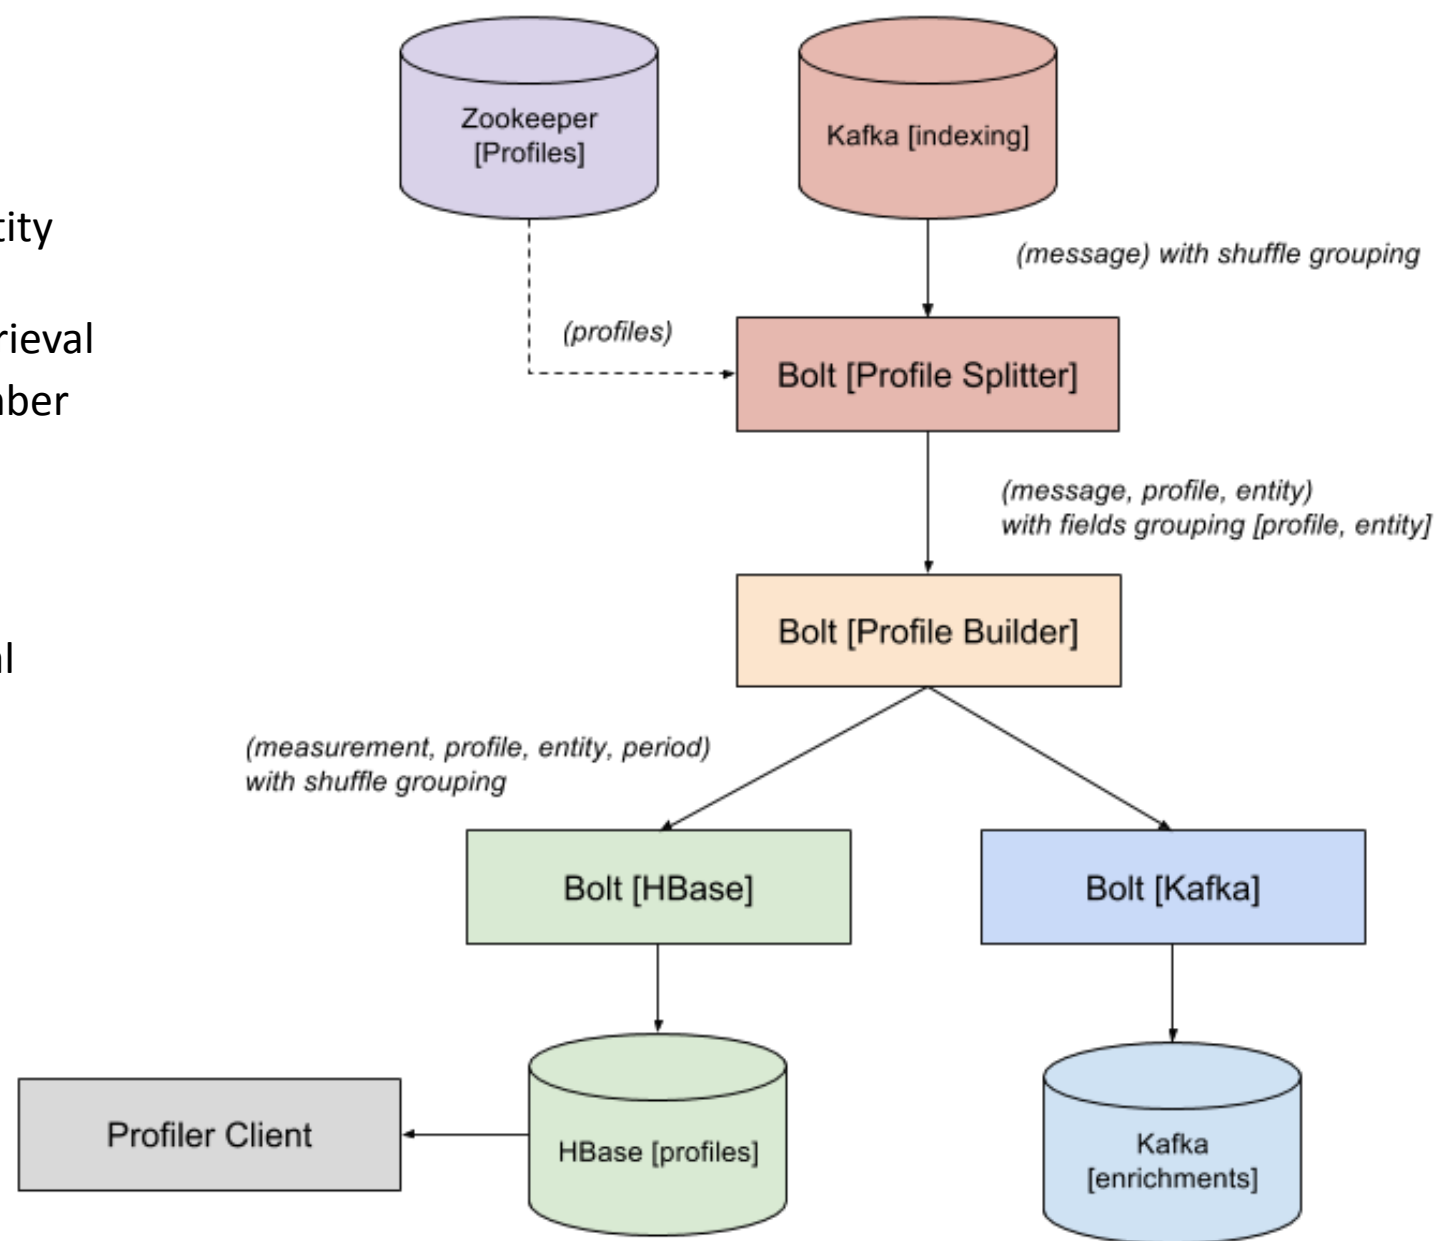
- Persists profile state in HBase
- Generates a row key using a salt, profile name, entity and time
- Row key needs to be deterministic to allow for retrieval
- The time period is a monotonically increasing number identifying each period since the epoch

## ■ Kafka Bolt

- Pushes profile data back into Metron
- Allows the user to create alerts based on abnormal profile values

## ■ Profiler Client

- Based on input parameters, calculates all of the necessary row keys
- Submits a multi-get using these row keys
- Some row keys may 'hit' and others 'miss'



# Q&A

- ◆ Questions?
- ◆ Join the community
  - ◆ <http://metron.apache.org/>
  - ◆ <https://github.com/apache/metron>
- ◆ More information on the Profiler
  - ◆ <http://metron.apache.org/current-book/metron-analytics/metron-profiler>